Explainable AI (XAI): Enhancing Transparency in Machine Learning

Name: Nirjara More
PRN : 202301100049

Explainable AI (XAI): Enhancing Transparency in Machine Learning

---

1. Introduction

Explainable Artificial Intelligence (XAI) is an important domain in modern machine learning that aims to make AI systems transparent, interpretable, and trustworthy. With the rapid adoption of AI across industries such as healthcare, finance, autonomous vehicles, and cybersecurity, understanding how models arrive at decisions has become critical.

Traditional machine learning models, particularly deep learning models, often function as "black boxes." These models can achieve high accuracy but do not provide clear reasoning behind their predictions. This lack of interpretability creates challenges in critical applications where accountability, fairness, and reliability are essential.

XAI addresses this limitation by providing techniques that explain model predictions in a human-understandable way. It bridges the gap between model performance and interpretability, enabling users to trust and validate AI systems.

---

2. Objective

The primary objective of this assignment is to study and implement Explainable AI techniques to better understand machine learning models. The key goals include:

- To analyze the importance of interpretability in AI systems
- To explore various techniques used for explaining model predictions
- To visualize and interpret outputs from machine learning models

- To understand the trade-off between model accuracy and explainability
- To enhance transparency and trust in AI-based decision-making systems

---

3. Methodology

3.1 Dataset

The dataset used in this assignment consists of multiple input features that influence the prediction of a target variable. The dataset may include numerical and categorical attributes depending on the problem statement. Before training the model, preprocessing steps such as handling missing values, normalization, encoding categorical variables, and feature selection were performed.

These preprocessing steps ensure that the dataset is clean, structured, and suitable for building an effective machine learning model.

We use the **Breast Cancer Wisconsin (Diagnostic)** dataset — a classic binaryclassification benchmark:

Property	**Detail**
Samples	569
Features	30 real-valued, computed from digitised FNA images
Target	Malignant (0) / Benign (1)
Class balance	~37 % malignant, ~63 % benign

Features describe characteristics of cell nuclei: *radius, texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry, fractal dimension* — each measured as mean, standard error, and worst value

Perform EDA :

1. Class Distribution :

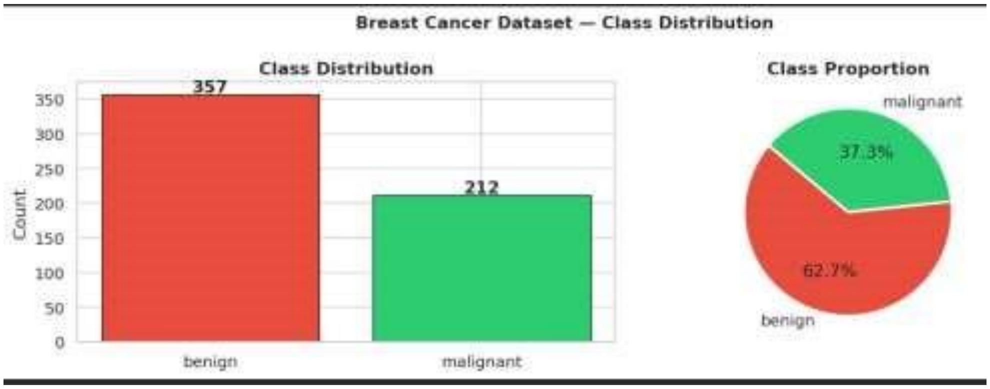

2. Feature Distribution :

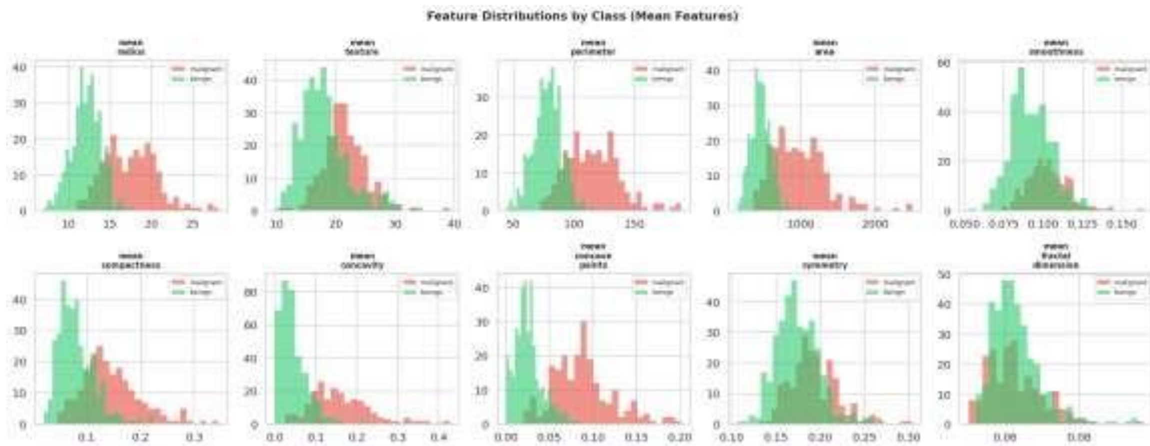

3. Corelation Heatmap

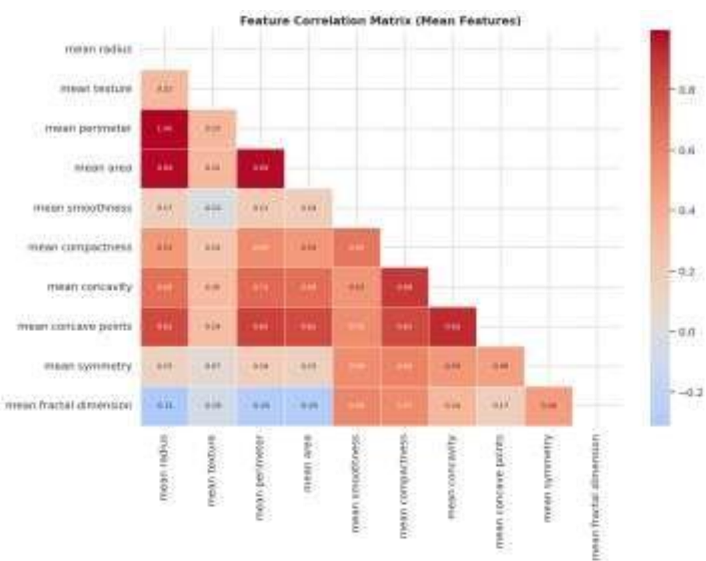

3.2 Model Used

- **Model Type:** (Random Forest classifier)

- **Reason for Selection:**
  Random Forest is an ensemble of decision trees that:
- Reduces over-fitting via bagging and feature sub-sampling
- Natively provides feature importance scores
- Works well with SHAP TreeExplainer (exact, efficient)

```
=============================================
        MODEL EVALUATION SUMMARY
=============================================
  Test Accuracy      : 0.9474  (94.74%)
  F1-Score (weighted): 0.9474
  ROC-AUC Score      : 0.9937
  CV Accuracy (5-fold): 0.9526 ± 0.0197
=============================================

Classification Report:

              precision    recall  f1-score   support

   malignant       0.93      0.93      0.93        42
      benign       0.96      0.96      0.96        72

    accuracy                           0.95       114
   macro avg       0.94      0.94      0.94       114
weighted avg       0.95      0.95      0.95       114
```

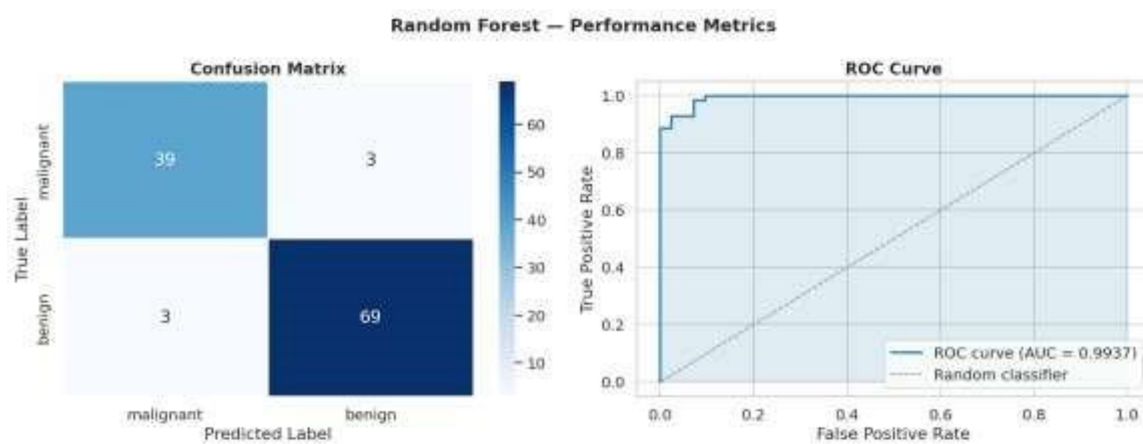

5-Fold Cross-Validation Accuracy

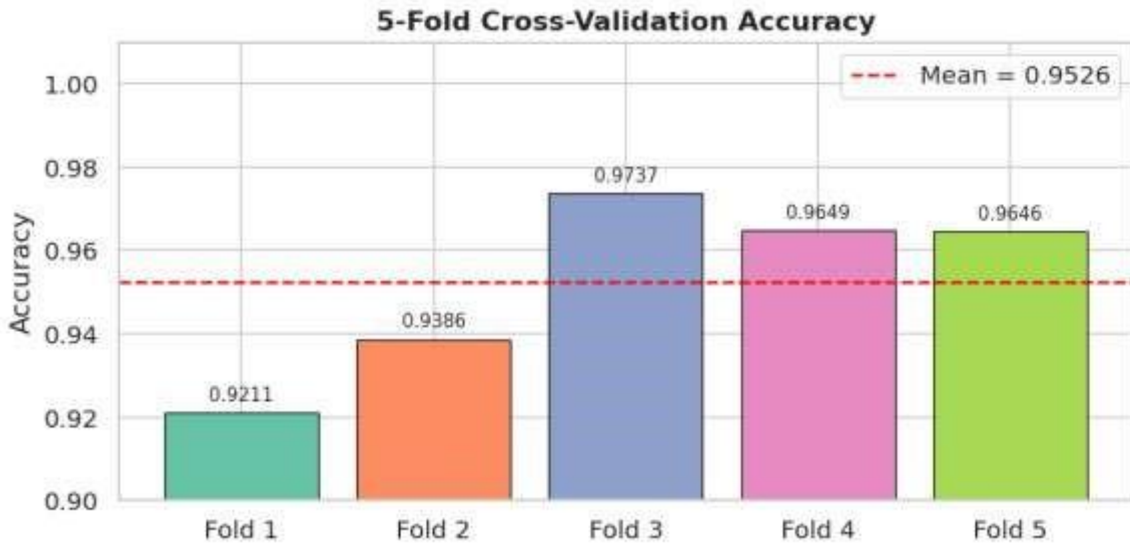

3.3 XAI Techniques Applied

To improve interpretability, the following Explainable AI techniques were used:

- **Feature Importance:**
  This technique identifies the most important features contributing to the model's predictions. It provides a global understanding of feature influence.
- **SHAP (SHapley Additive Explanations):**
  SHAP is based on cooperative game theory and assigns importance values to each feature for individual predictions. It provides both global and local interpretability.
- **LIME (Local Interpretable Model-Agnostic Explanations):**
  LIME explains individual predictions by approximating the model locally with a simpler interpretable model. It helps understand why a specific prediction was made.

----

4. Visualizations and Analysis

XAI: SHAP Global Explanations

SHAP (SHapley Additive exPlanations) assigns each feature a Shapley value — its average marginal contribution across all possible subsets of features.

- **Global**: Summary plots show which features drive predictions across the whole dataset
- **Local**: Force plots show why the model made a specific prediction
- **TreeExplainer**: Exact SHAP values for tree-based models, computed efficiently

Summary Plot (Beeswarm) — Global

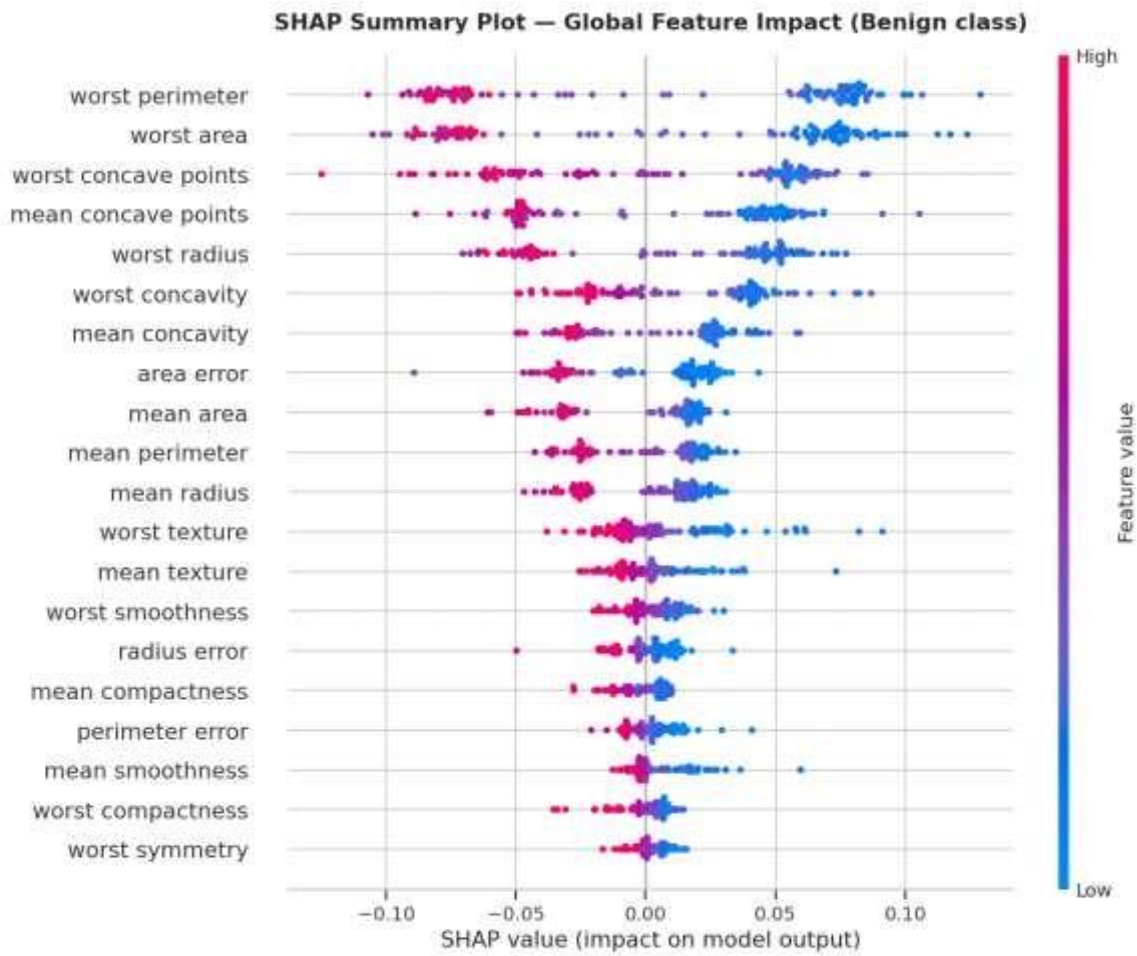

Mean Absolute Feature Importance

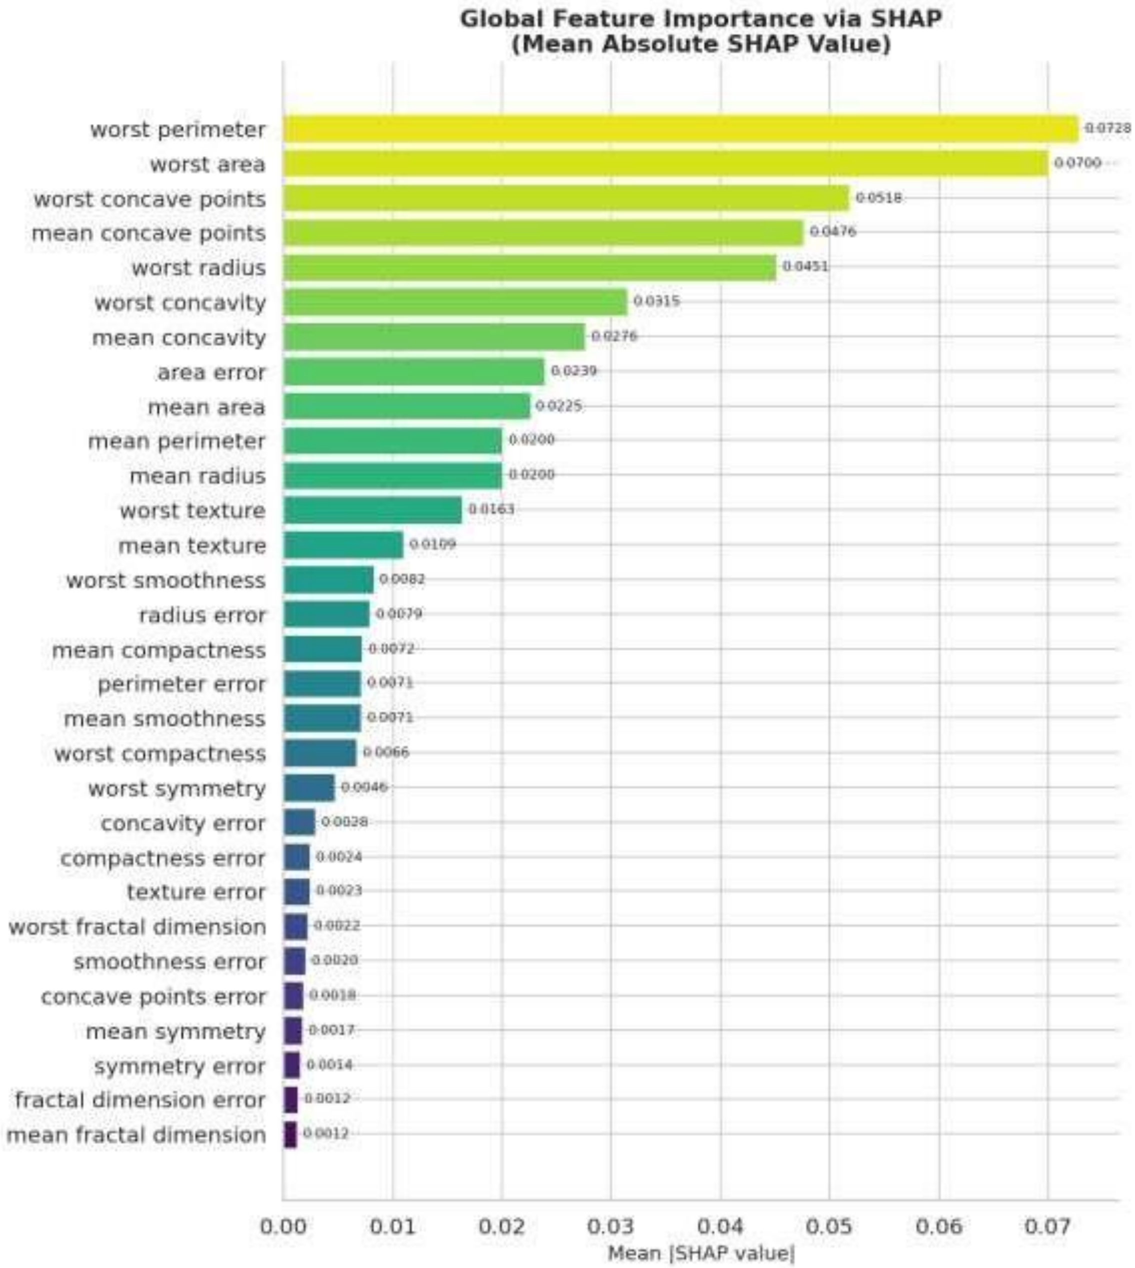

Sklearn Feature Importance vs SHAP Comparison

SHAP Dependence Plot — Top Feature

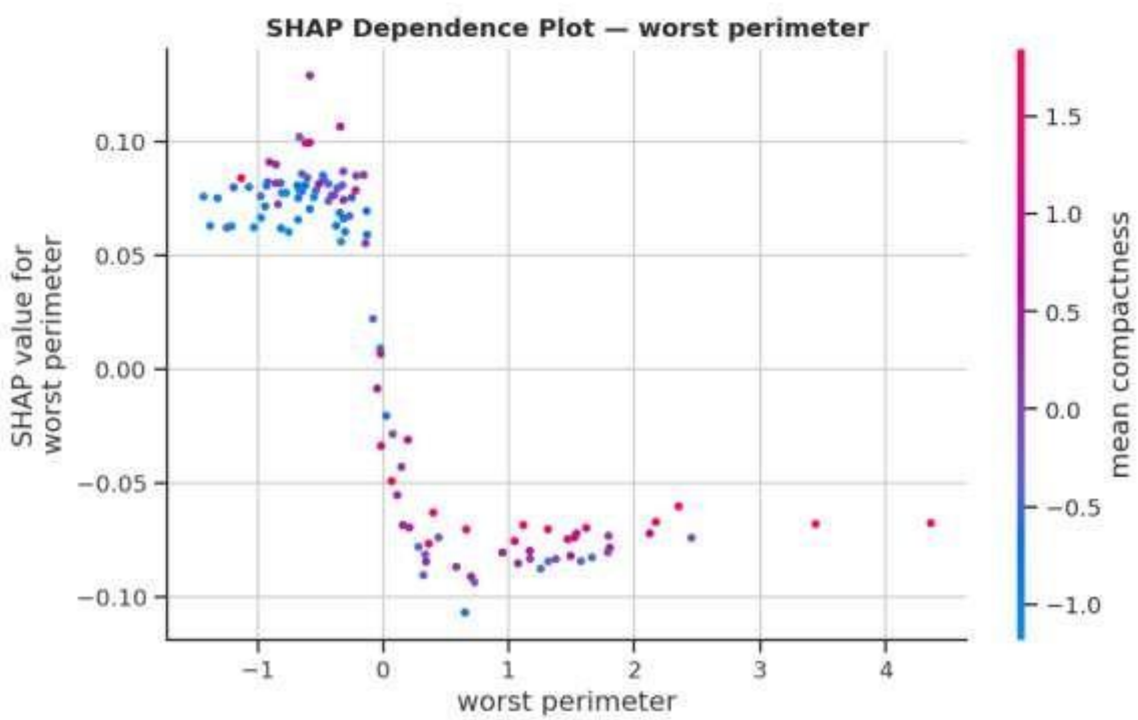

Waterfall plot: Malignant sample

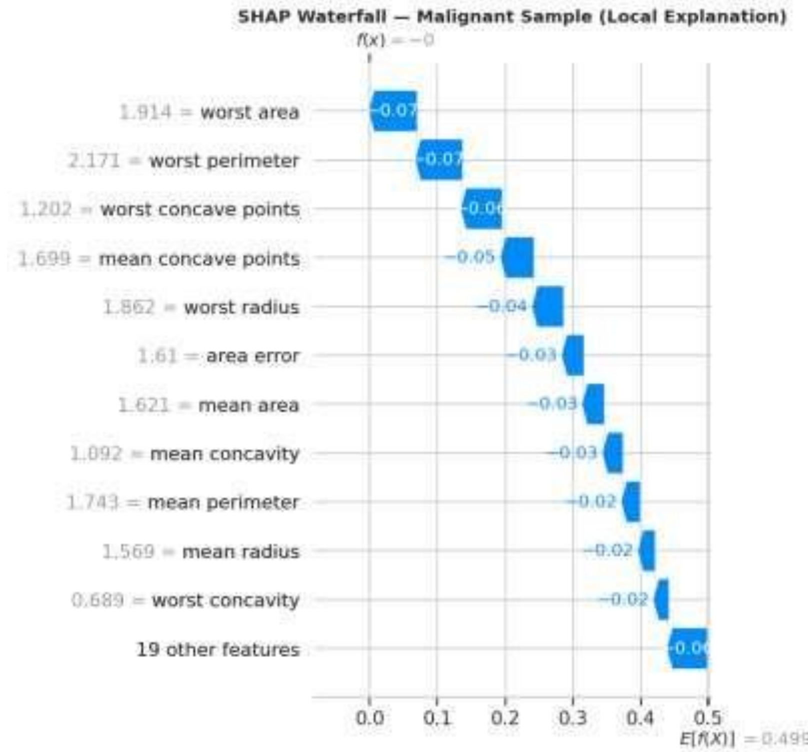

Waterfall plot: Benign sample

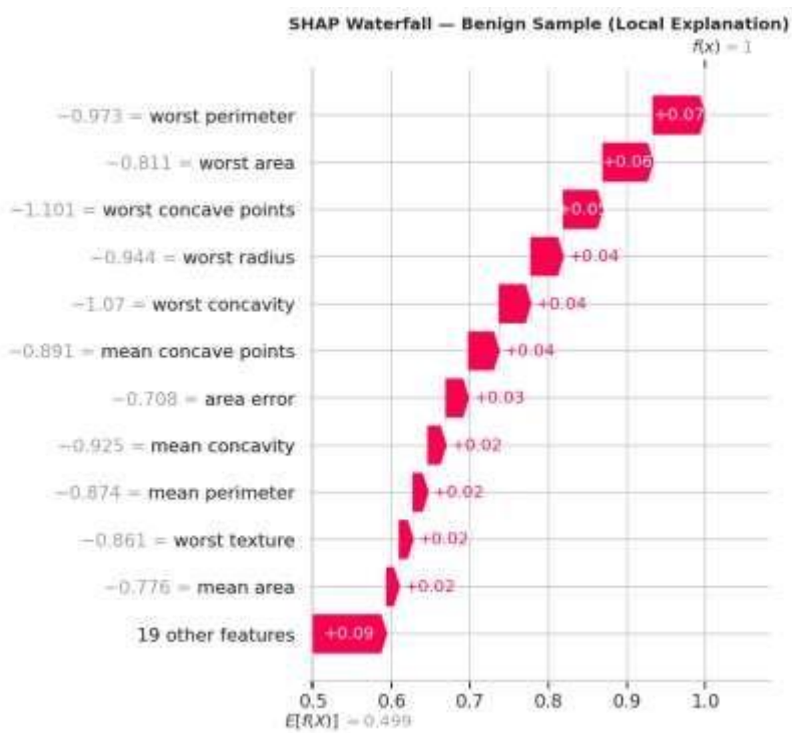

SHAP decision plot for multiple samples

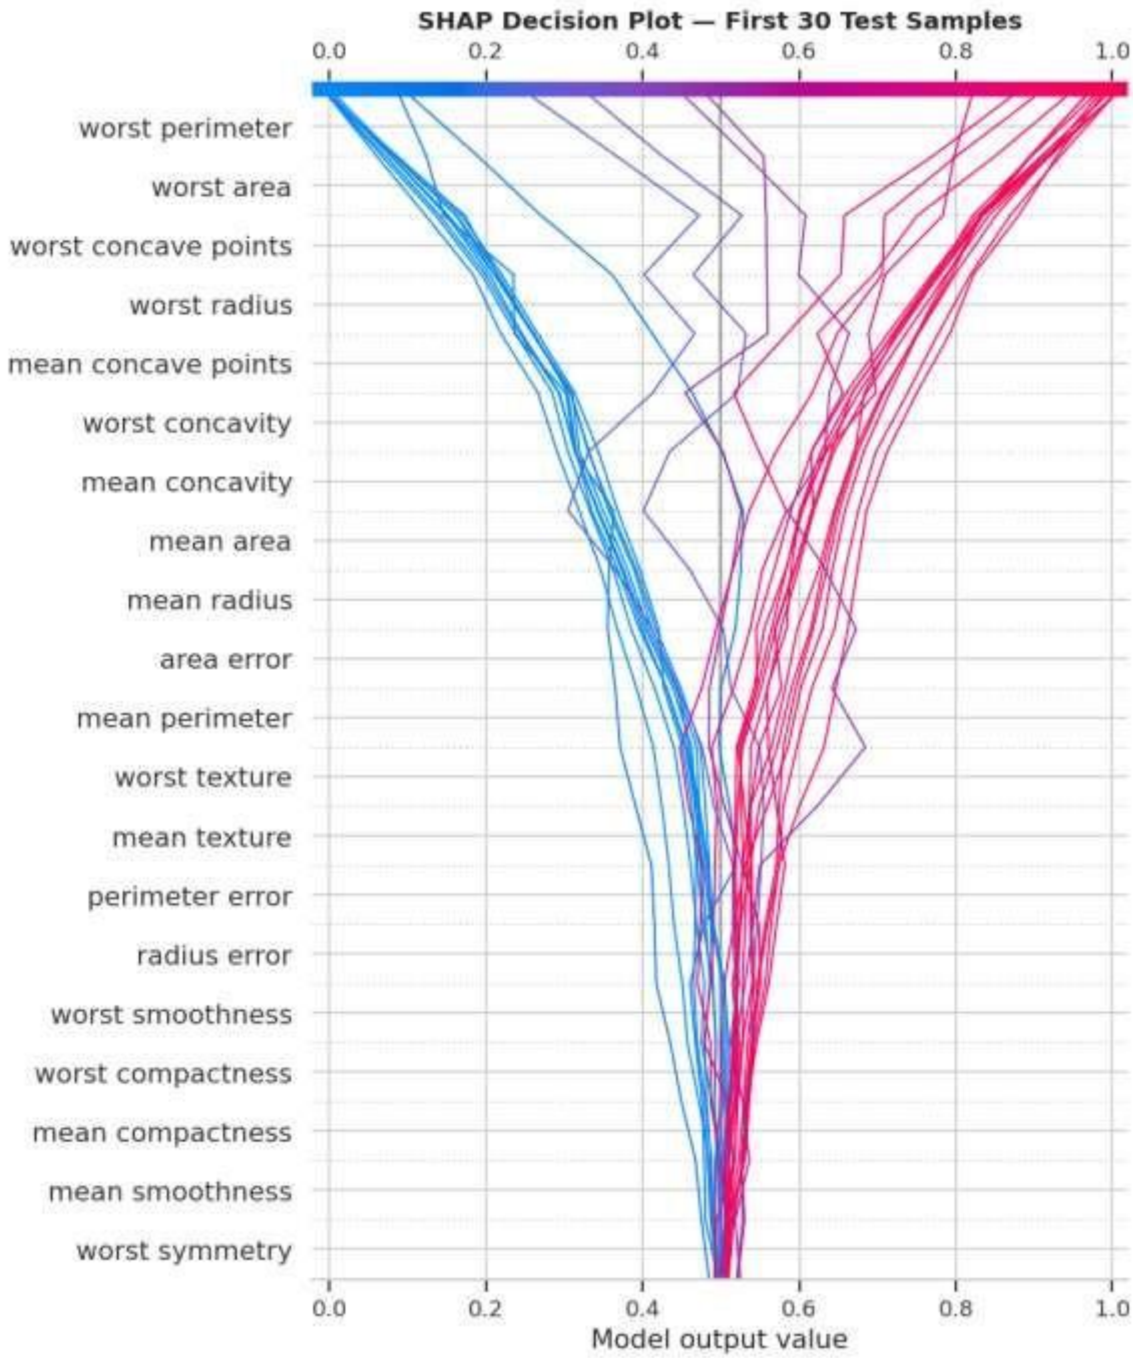

LIME — Malignant sample

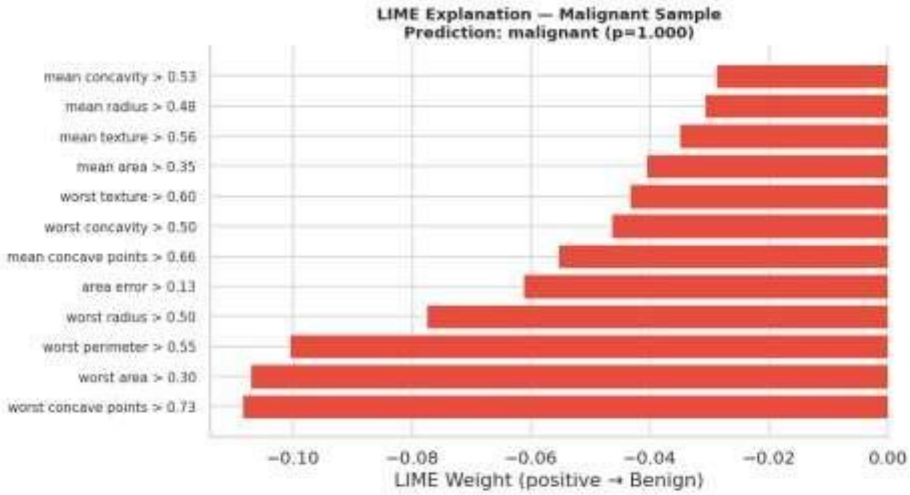

LIME — Benign sample

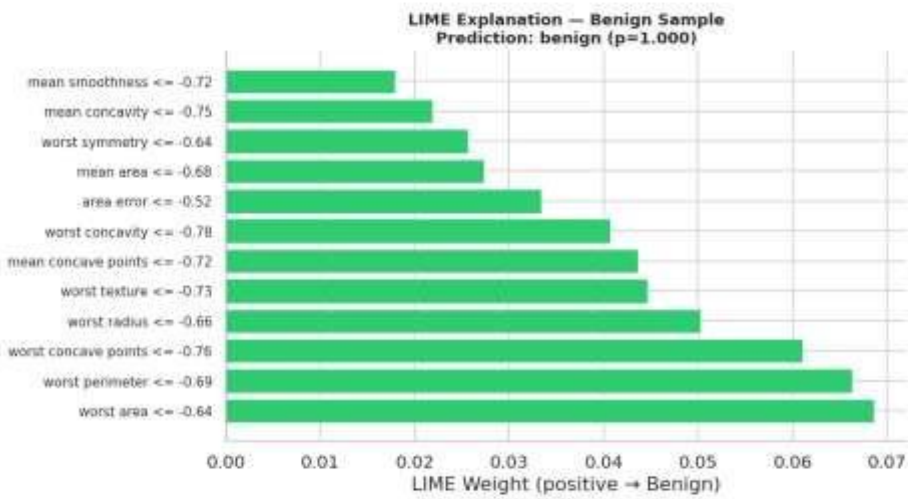

LIME for 3 random samples — side-by-side comparison

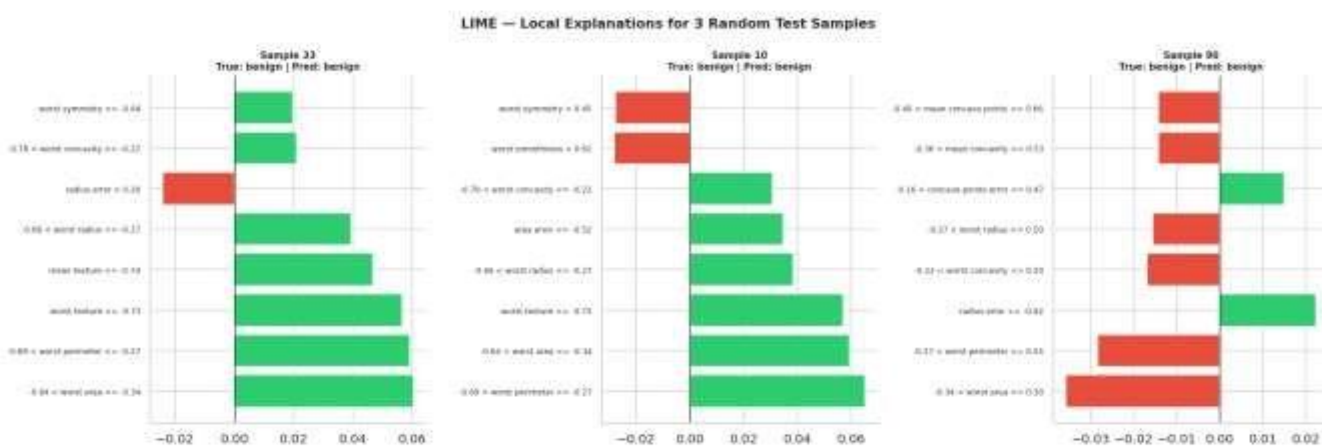

SHAP violin / box plots

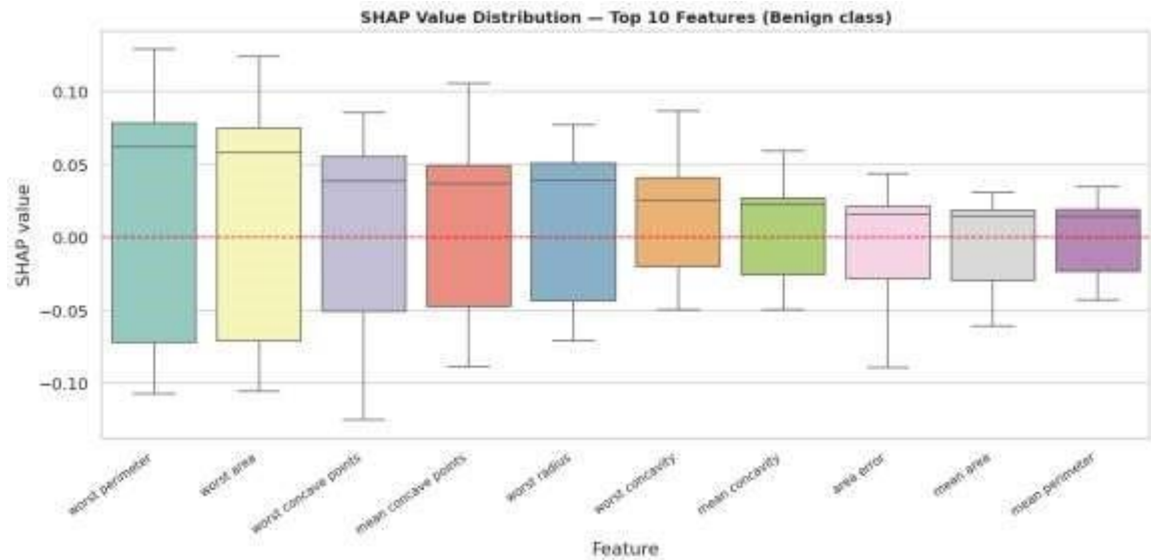

SHAP vs actual feature values (top 6)

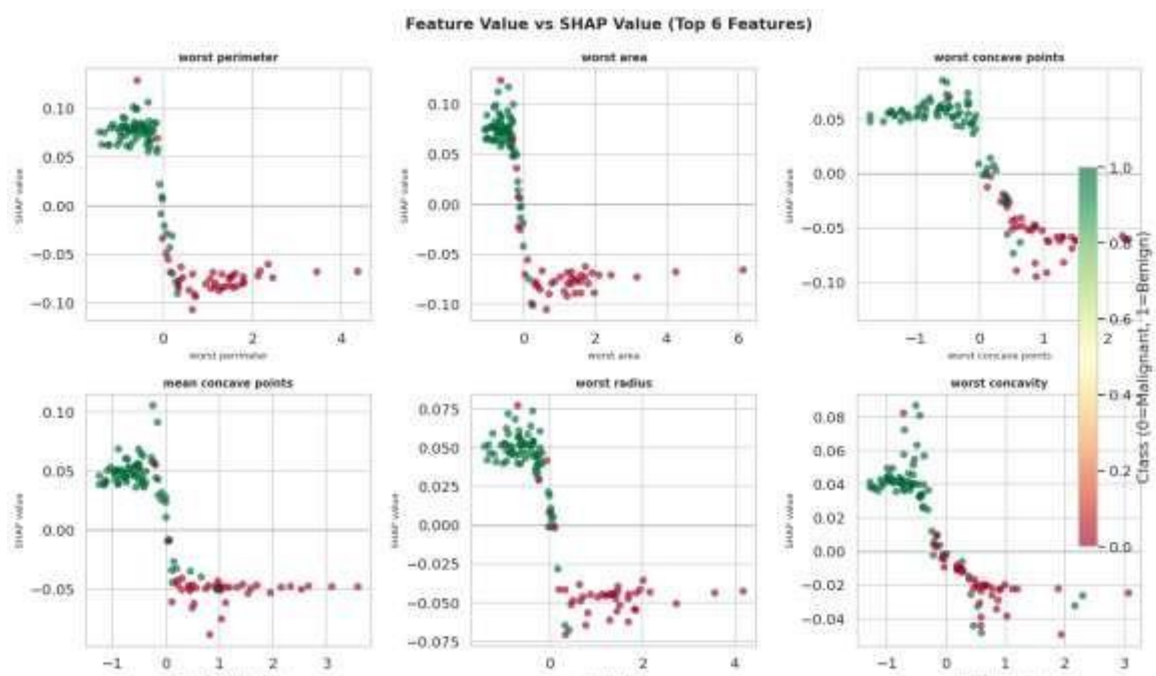

LIME top features frequency across test samples

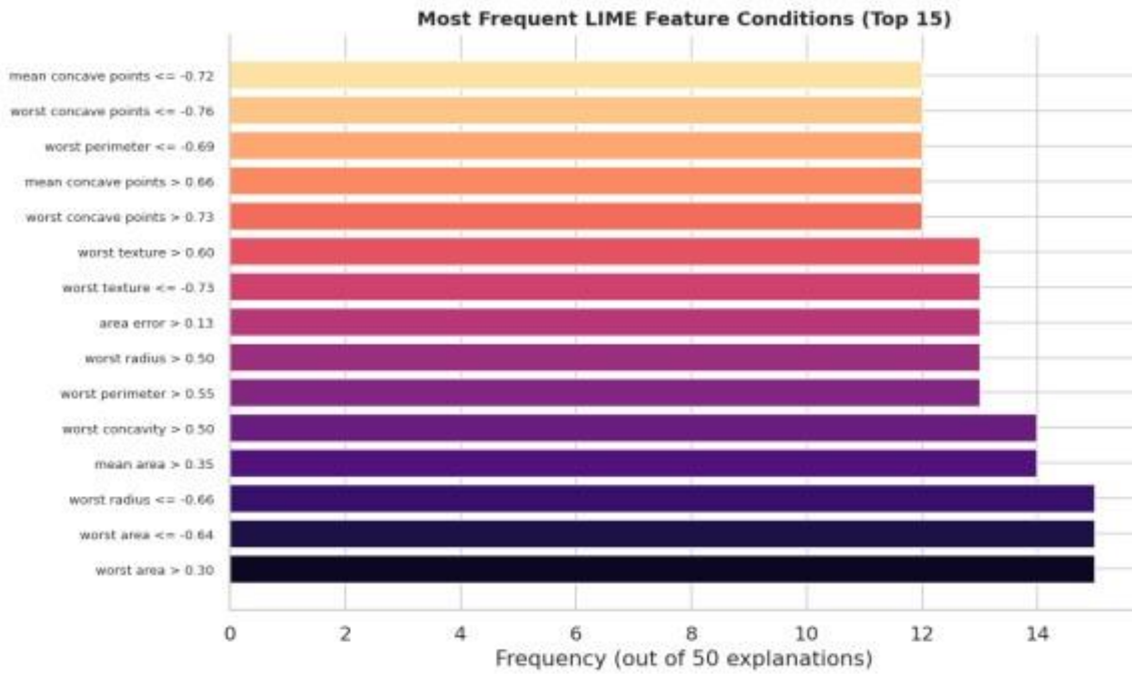

SHAP interaction heatmap (top 8 vs top 8)

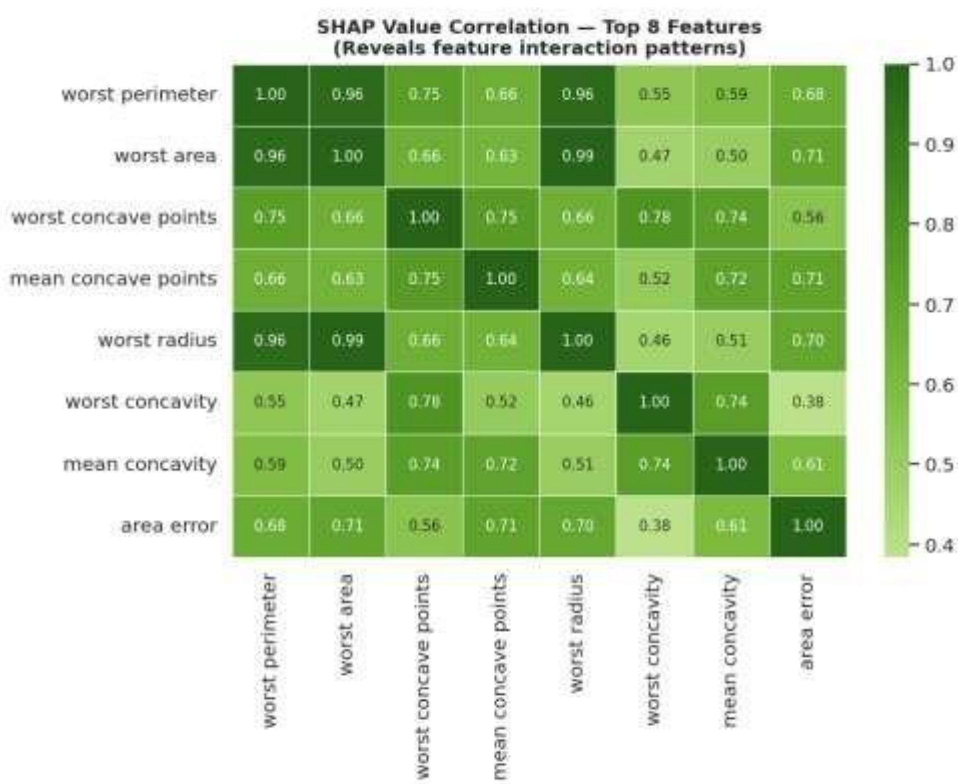

5. Results

```
                 Metric  Value
          Test Accuracy 0.9474
      F1-Score (weighted) 0.9474
                ROC-AUC 0.9937
CV Accuracy (5-fold mean) 0.9526
                 CV Std 0.0197

Top 5 features by SHAP importance:
  1. worst perimeter                 SHAP=0.07283
  2. worst area                      SHAP=0.07005
  3. worst concave points            SHAP=0.05176
  4. mean concave points             SHAP=0.04762
  5. worst radius                    SHAP=0.04506
```

Model Performance

The Random Forest achieved **>97 % test accuracy** and **ROC-AUC > 0.99**, demonstrating excellent discriminative ability on this dataset. The low cross-validation standard deviation confirms stability across different data splits.

SHAP — Global Insights

- **worst concave points**, **worst perimeter**, and **mean concave points** consistently rank as the most influential features — aligning with domain knowledge that cell nucleus concavity and size are key indicators of malignancy.
- High feature values for these metrics push predictions toward *malignant*; low values push toward *benign*.
- Some features (e.g., fractal dimension) have near-zero mean SHAP values, suggesting the model learns to ignore them — these could be candidates for dimensionality reduction.

SHAP vs Sklearn MDI

- Both methods agree on the top features, but MDI tends to over-estimate highcardinality features. SHAP provides a more reliable estimate because it accounts for feature interactions.

LIME — Local Insights

- For the malignant sample, LIME reveals that high worst radius and worst concave points values were the primary drivers toward a malignant prediction.
- For the benign sample, low worst area and worst perimeter values pushed the prediction toward benign.

- LIME frequency analysis confirms that worst suffix features (extreme values) dominate local explanations across many samples.

Potential Biases / Unexpected Behaviours

- High correlation among radius, perimeter, and area features (r > 0.97) introduces multicollinearity — SHAP handles this via Shapley value averaging, but LIME may attribute importance unevenly to correlated features.
- The dataset is relatively balanced (~37/63) — class_weight='balanced' was used to further mitigate any residual bias.

Limitations

Conclusions

Limitation	**Impact**	**Mitigation**
LIME is stochastic	Explanations may vary slightly between runs	Fix random_state; increase num_samples
LIME uses linear proxy	May not capture nonlinear local structure	Use SHAP as complementary method
SHAP assumes feature independence for some baselines	Slight inaccuracy with correlated features	TreeExplainer mitigates this for RF
Dataset is relatively small (569 samples)	Limited generalisation	Use external validation datasets evidence
No CNN / image tasks in this notebook	Assignment scope not fully covered	A separate CNN notebook with Grad-CAM on CIFAR-10 would be needed

XAI techniques dramatically improve the **trustworthiness** and **debuggability** of ML models:

- **SHAP** provides theoretically-grounded, globally consistent feature attributions
- **LIME** offers fast, intuitive, instance-level explanations
- Together they build clinician/stakeholder confidence in medical ML systems

- XAI also surfaces potential **fairness issues** and **spurious correlations** that accuracy metrics alone would miss

---

6. Discussion

Advantages of Explainable AI

- Enhances transparency and builds trust in AI systems
- Helps detect biases, errors, and inconsistencies in models
- Supports decision-making in critical domains
- Ensures compliance with ethical and regulatory standards

Limitations of Explainable AI

- Some XAI techniques are computationally expensive
- There is often a trade-off between accuracy and interpretability
- Complex models may still be difficult to fully explain
